

AI-Native SDLC with Matt Pocock Skills

Oct 2, 2026 · @Dorin

Summary

This SDLC uses Anthropic's six-stage loop (Plan, Design, Build, Test, Deploy, Maintain) as the skeleton and Matt Pocock's skills as the working method inside each stage. Anthropic supplies the governance: a versioned artifact per stage, hooks that block what must never happen, evals, and humans at every approval gate. Pocock supplies the engineering discipline: grilling before building, tracer-bullet tickets, TDD at agreed seams, two-axis code review, and regular architecture cleanup.

Five rules hold the whole thing together:

1. **Every stage ends in a committed artifact.** intent.md, spec issue, tickets, diff with tests, merged PR. The artifact chain is the audit trail.
2. **Humans decide, Claude executes.** A person approves intent, spec, plan, merge and release; Claude drafts, builds, verifies and fixes.
3. **Skills advise, hooks enforce.** Pocock's skills shape how Claude works. Anything that must hold without exception is backed by a hook or branch protection.
4. **Claude verifies its own work before a human sees it.** Typecheck, tests, full suite and /code-review run before a PR is opened.
5. **The loop closes.** Incidents and repeated review mistakes become new intents, evals and CLAUDE.md rules.

What each source contributes

The two sources barely overlap: Anthropic defines the stages, artifacts and gates, and Pocock defines how the work inside a stage is done.

	Anthropic AI-Native SDLC Playbook	mattpocock/skills
Answers	What happens in which stage, who approves, how it is enforced	How Claude should think and work while doing a task
Core unit	Stage plus versioned artifact (intent.md, spec.md, plan.md)	Skill invoked by the user or the model
Control	Hooks, managed settings, branch protection, REVIEW.md, evals in CI	Prompts that impose discipline (grilling, TDD, two-axis review)

	Anthropic AI-Native SDLC Playbook	mattpocock/skills
Human role	Owens intent, approves gates, governs the loop	Answers the grilling, picks tickets, reviews output
Gap it leaves	Says little about engineering craft inside Build	No gates, no enforcement, no deploy or ops stage


Anthropic's [security write-up](#) adds risk tiering of the codebase, /security-review during coding, sandboxed agents with egress allowlists, and logging of every automated approval. Pocock's skills are installed with `npx skills@latest add mattpocock/skills` followed by a one-time `/setup-matt-pocock-skills` per repo, which configures the issue tracker, triage labels and doc locations.

The six stages

Each stage starts from the previous stage's artifact and ends at a gate a human owns. Claude does the drafting and execution in every stage; the human's job is to sharpen intent and say yes or no at the gate.

Every stage ends in an artifact and a human gate

Each box: artifact, main Pocock skills, exit gate



AI-native SDLC loop · 6 stages, 6 gates

Work flows clockwise; Maintain feeds new intents back into Plan, so the loop never ends at deploy.

Stage	Human does	Claude does	Pocock skills	Artifact	Exit gate
1. Plan	Writes the ask in their own words; decides if it is worth doing	Triages inbound issues, interviews the requester until the intent is unambiguous	triage, grill-me, wait-what, to-questionnaire, then to-intent (team skill) to save the result	specs/<slug>/intent.md	G1 Product owner approves intent
2. Design	Answers the grilling; settles trade-offs; policy owners review risky areas	Builds the domain model, researches, prototypes the open questions, writes the spec	grill-with-docs, domain-modeling, research, prototype, codebase-design, to-spec	Spec issue in the tracker, linked from intent.md; updated GLOSSARY.md	G2 Owner approves spec; security or compliance owner signs if risk tier is high
3. Build	Corrects the ticket plan before code; runs 2 to 3 sessions in parallel	Splits the spec into tracer-bullet tickets, then implements one ticket per session with TDD	to-tickets (or wayfinder for multi-week work), implement, tdd, diagnosing-bugs, resolving-merge-conflicts, handoff	Tickets with blocking edges (this is Anthropic's plan.md), then commits	G3 Engineer approves tickets before the first line of code
4. Test	Reviews only what survived automated checks	Runs typecheck, single tests, full suite, two-axis /code-review, /security-review, evals	tdd, code-review, then review-report (team skill) to save the result	Test files in the diff, CI test and coverage reports, review report comment on the PR	G4 CI green, evals pass, no open Important findings
5. Deploy	Code owner approves the merge; release owner approves production	Opens the PR, addresses review comments, writes wizards for manual steps	code-review, wizard, then to-release (team skill) for the release record	Merged PR with review report and approval; draft GitHub Release from to-release; deployment record from CI	G5 Human merge approval (the authoring agent can never approve). G6 Release approval for production

Stage	Human does	Claude does	Pocock skills	Artifact	Exit gate
6. Maintain	On-call triages findings; decides which become work	Diagnoses incidents, files new intents, scans for architecture decay every few days	diagnosing-bugs, improve-codebase-architecture, triage	New <code>intent.md</code> , new evals, CLAUDE.md rules	Loops back to G1

Two Pocock skills sit outside the stages. **handoff** compacts a session into a document so the next session can continue a ticket. **ask-matt** routes to the right skill when nobody remembers which one fits.

How the Plan and Design artifacts get written

Neither grill-me nor to-spec writes a repo file. grill-me only runs the interview, and to-spec publishes the spec as an issue in your tracker with the ready-for-agent label. So the two artifacts are produced like this:

1. **intent.md**: after grill-me ends, run **to-intent**, a small team skill you add next to Pocock's. Like to-spec, it does not interview; it condenses the conversation into `specs/<slug>/intent.md` and opens a PR. Merging that PR is G1.
2. **Spec**: run grill-with-docs, then to-spec. The spec lives in the tracker, and the intent.md PR links to it. The product owner applying ready-for-agent is G2.

Keeping the spec in the tracker follows the playbook's advice to name one system of record per artifact, and it is where to-tickets and code-review already look.

How the Test artifacts get written

Pocock's tdd writes test files, but code-review and /security-review only report in chat. Stage 4 therefore keeps three artifacts:

1. **Tests**: the test files tdd writes are committed in the diff. The test plan they follow is the Testing Decisions section of the spec issue that to-spec writes.
2. **Test run**: CI saves a JUnit XML and coverage report with every build, and the PR links to them.
3. **Review report**: run **review-report**, a second team skill. It runs typecheck, the full suite, /code-review and /security-review, then posts one PR comment with the G4 checklist, the tests added, findings ranked Important or Nit, and a verdict. A rerun edits the same comment, and the skill never approves the PR.

The report is a PR comment rather than a repo file, so it sits next to the diff it judges and adds no file to every PR.

How the Deploy artifacts get written

Deploy has two gates, and each leaves a record:

1. **G5 merge:** the merged PR is the artifact. It already carries the review report, the human approval and the CI results.
2. **G6 release:** run **to-release**, a third team skill. It drafts a release record for each production deploy and saves it as a draft GitHub Release. The record lists every PR that ships with its intent, spec, risk tier, approver and review verdict. It also flags PRs missing an approval or a Ready report, lists risky changes with a rollback for each, and links a wizard script for manual steps.
3. **Deploy log:** CI deploys through the production GitHub environment, which records each deploy against its commit and waits for the release owner's approval.

The release owner reads the draft, approves the production environment, then publishes the release. The skill never deploys, publishes or approves anything.

Skill-to-stage map

Every skill in the repo has a home; only teach is optional. "User" skills are run as slash commands; "Model" skills Claude picks up on its own.

Skill	Stage	Invoked by	Use it when
setup-matt-pocock-skills	Setup	User	Once per repo, to point skills at the tracker, labels and docs folders
ask-matt	Any	User	You are unsure which skill fits
triage	1 Plan, 6 Maintain	User	Moving inbound issues through triage states
grill-me / grilling	1 Plan	User / Model	The ask is vague; Claude interviews until every branch is resolved. Writes no file, so run to-intent next
wait-what	1 Plan	User	A stakeholder message is unclear; re-pitch it in plain words
to-questionnaire	1 Plan, 2 Design	User	A decision needs async input from people not in the session
grill-with-docs	2 Design	User	Aligning on requirements while updating GLOSSARY.md

Skill	Stage	Invoked by	Use it when
domain-modeling	2 Design	Model	Terms are fuzzy or overloaded
research	2 Design	Model	A question needs primary sources and a cited answer
prototype	2 Design	Model	A design question is cheaper to answer with throwaway code
codebase-design	2 Design, 3 Build	Model	Choosing seams and module boundaries; these become the TDD seams
to-spec	2 Design	User	Turning the design discussion into a spec issue in the tracker, labelled ready-for-agent
to-tickets	3 Build	User	Turning the spec issue into tracer-bullet tickets with blocking edges
wayfinder	3 Build	User	Work spans many sessions and needs a decision map
implement	3 Build	User	Building one ticket; it drives tdd and code-review
tdd	3 Build, 4 Test	Model	Writing code at an agreed seam, red-green-refactor
diagnosing-bugs	3 Build, 6 Maintain	Model	Anything is broken: reproduce, minimise, hypothesise, instrument, fix
resolving-merge-conflicts	3 Build	Model	A branch conflicts with main
handoff	3 Build	User	Ending a session mid-ticket
code-review	4 Test, 5 Deploy	Model	Before every PR: Standards axis and Spec axis. Output stays in chat, so review-report runs it and saves the result
wizard	5 Deploy	Model	A release has manual steps a human must run; to-release calls it
improve-codebase-architecture	6 Maintain	User	Every few days, to catch entropy before it compounds

Skill	Stage	Invoked by	Use it when
writing-for-agents	Setup	Model	Writing CLAUDE.md, REVIEW.md and skills
teach	Optional	User	Onboarding a person onto the codebase

Roles for a 10-developer team

With 10 developers, every gate keeps a separate human owner, but most roles are hats the developers wear on a rotation rather than new hires. The one structural rule: the person who drove Claude on a change never gives that change its G5 approval.

Role	Held by	Gates	Notes
Product owner	Product manager, or the tech lead if there is none	G1, G2	Approves intent.md and the spec issue for every feature
Feature driver	Any developer, one per feature	G3	Runs the grilling, approves tickets, steers 2 to 3 Claude sessions
Code owners	2 to 3 developers per area in CODEOWNERS	G5	Approve merges for areas they do not author in that change
Policy owners	One developer each for security and for data or compliance	G2 on high-risk specs	Also own the matching policy skill and hooks
Release owner	Weekly rotation	G6	Approves production deploys and runs any wizard steps
On-call	Weekly rotation, can be the same person as release owner	Maintain	Triages monitoring findings and turns the real ones into intent.md
SDLC steward	One developer, rotating monthly	None	Owens CLAUDE.md, REVIEW.md, evals and hooks; reviews eval pass rates and samples automated approvals

Plan for review becoming the bottleneck. Ten developers running 2 to 3 sessions each can open 20 or more PRs a day, so keep PRs to one tracer-bullet ticket each, let Claude's review clear the Nits before a human looks, and reserve human review time each day rather than reviewing between tasks.

Gates and controls

A gate passes only on a human decision or a deterministic check, never on a skill's say-so. Skills tell Claude how to work well; hooks, branch protection and CI make the non-negotiables hold.

Gate	Passes when	Who decides	Enforced by
G1 Intent	intent.md states the problem, who has it and what done looks like, in the requester's words	Product owner	PR to <code>specs/</code> with CODEOWNERS on that folder
G2 Spec	The spec issue is approved; GLOSSARY.md is updated; high-risk areas signed by their policy owner	Product owner, plus security or compliance owner by risk tier	Only the product owner may apply the ready-for-agent label; risk tier label
G3 Plan	Tickets are thin vertical slices with blocking edges and named TDD seams	Engineer	Tickets moved to Ready in the tracker
G4 Verified	Typecheck, full suite and evals green; /code-review and /security-review have no open Important findings; the review report says Ready for human review	CI	Required status checks
G5 Merge	A human code owner approves; Claude's review informs but never approves	Code owner	Branch protection, required reviewers

Gate	Passes when	Who decides	Enforced by
G6 Release	Release owner has checked the to-release draft and approves the production environment	Release owner	Deploy hook or environment protection rule

Controls that run all the time:

- **Hooks** block edits to protected paths (migrations, generated code, secrets, CI config), run the formatter after every edit, and stop pushes to main.
- **Managed settings** fix permissions and sandboxing so an engineer cannot loosen them per session. Agents run with an egress allowlist.
- **Risk tiers** decide how much autonomy each area gets: dev deploys autonomously, staging is mixed, production always needs G6. Auth, payments and data-deletion paths always need a human reviewer at G5.
- **REVIEW.md** defines what Claude's PR review checks for and how it ranks severity (Important vs Nit). Pocock's code-review reads `CODING_STANDARDS.md` for its Standards axis and the issue referenced in commit messages (or a path you give it) for its Spec axis, so keep standards in the repo and reference the spec issue in every commit.
- **Evals**: 20 to 50 real past tasks, rerun in CI whenever CLAUDE.md, a skill or a hook changes. A config change that drops the pass rate does not merge.
- **Audit**: the artifact chain plus logged automated approvals; sample a few automated approvals each week.

Repo setup

One repo layout serves both sources: the artifacts Anthropic asks for live where Pocock's skills already look for them.

```

CLAUDE.md                # conventions, commands, mistakes Claude made
twice
GLOSSARY.md              # domain language, kept current by grill-with-docs
CODING_STANDARDS.md      # Standards axis for /code-review
REVIEW.md                # what Claude's PR review checks and how it ranks
severity
specs/<slug>/intent.md    # Stage 1, written by to-intent, approved at G1
evals/                   # 20 to 50 real past tasks
.claude/settings.json    # hooks and permissions
.claude/skills/          # Pocock skills plus to-intent, review-report, to-
release and your policy skills
.github/CODEOWNERS       # owners for specs/, auth, payments, infra

```

The spec is not a repo file: to-spec publishes it as a tracker issue, and to-tickets and code-review read it from there.

Setup checklist:

- ☐ Install the skills: `npx skills@latest add mattpocock/skills`, then run `/setup-matt-pocock-skills`
- ☐ Add the to-intent team skill to `.claude/skills/to-intent/` (intent files go in `specs/`); the tracker configured by `/setup-matt-pocock-skills` holds specs and tickets
- ☐ Write CLAUDE.md with build, test and typecheck commands, using writing-for-agents. Add this line so Pocock's skills use our glossary name: "The domain glossary is GLOSSARY.md. Wherever a skill refers to CONTEXT.md, read and update GLOSSARY.md instead." Pocock's skills hard-code CONTEXT.md, so check after the first grill-with-docs run that no CONTEXT.md was created
- ☐ Write CODING_STANDARDS.md and REVIEW.md
- ☐ Add an intent.md template under `specs/_template/`
- ☐ Add hooks: block protected paths, format on edit, block pushes to main
- ☐ Turn on branch protection with required CI checks and a required code owner review
- ☐ Add the review-report team skill, Claude PR review in CI (Claude Code GitHub Action or Code Review), and JUnit XML plus coverage uploads as CI build artifacts. Deploy through a GitHub environment named production with required reviewers, so every deploy is recorded and waits for G6
- ☐ Collect the first 20 eval tasks from recently merged PRs and run them in CI
- ☐ Label risk tiers in CODEOWNERS and the deploy pipeline

Rollout and metrics

Roll it out in Anthropic's four steps, starting wherever work stalls today: slow reviews mean start at Deploy, vague asks mean start at Plan.

1. **Manual loop.** Run the stages by hand on one real feature: grill-me, to-spec, to-tickets, implement. Save intent.md with to-intent and publish the spec issue with to-spec. Write CLAUDE.md and GLOSSARY.md as you go.
2. **Enforcement.** Add Claude PR review against REVIEW.md, the protected-path and production hooks, branch protection, and your first policy skill (for example a secure-API skill).
3. **Automation.** Let a committed intent.md or a spec issue labelled ready-for-agent kick off the next stage. Run 2 to 3 parallel sessions per engineer, one ticket each. Turn on evals in CI.
4. **Closed loop.** Monitoring opens intent.md files on its own, incidents become evals, and improve-codebase-architecture runs on a schedule.

Measure each stage with one leading and one lagging number, taken from the playbook:

Stage	Leading	Lagging
Plan	Hours from ask to committed intent.md	Share of approved intents that ship
Design	Hours from intent.md to approved spec issue	Rework after build starts
Build	First-pass merge rate	Diff still matches the tickets
Test	First-pass CI success rate	Defects that reach production
Deploy	Minutes to first review	Vulnerabilities found after merge vs before
Maintain	Minutes from incident to intent.md	Repeat incidents

Add one maintenance number the playbook leaves out: the share of merged PRs that are refactors from improve-codebase-architecture. If it falls toward zero, entropy is winning.